

WHITE PAPER | AUGUST 2020

MongoDB + Diamanti: A Hyper-Accelerated Story

AUTHORS:

Boris Kurktchiev, Diamanti

Jason Mimick, MongoDB

Contents

3	Overview
3	Goals of the Tests
3	Specifications
3	Hardware Used
4	Software Used
4	Test Results
8	Tests I - Workload A (50/50 read/write)
8	Tests III - Workload E (95/5 Scan/Update)
9	Findings
9	Test I - Workload A
10	Raw Numbers for Workload A:
11	Test II - Workload B
12	Raw Numbers for Workload B:
12	Test III - Workload E
14	Raw Numbers for Workload E:
15	Conclusion

Overview

Organizations today are embracing containers and Kubernetes for rapid development and portability of stateful applications. While containers and Kubernetes can abstract away much of the infrastructure, it is still an important consideration when containerizing high-performance databases such as MongoDB.

Enterprises need to ensure that the infrastructure can support multiple service level offerings (performance levels) and is self-managed across on-premises and cloud environments. In addition to this, these stateful applications contain business-critical information and demand advanced data services like data protection, monitoring, and fine-grained RBAC over a transactional data store with encryption at rest and in transit.

MongoDB has traditionally delivered excellent performance to value ratios. Typically, these deployments have been bare-metal or VM based. However, as cloud-native continues to grow and dominate enterprise software development and delivery practices, the ability to offer first-class enterprise-grade data services into containerized platforms such as Kubernetes is more urgent than ever before.

This document showcases the power of using MongoDB on Diamanti to drive I/O value as well as demonstrate the elasticity and usability driven by the adoption of Kubernetes and [MongoDB Enterprise Advanced](#).

Goals of the Tests

1. Showcase a resilient MongoDB architecture on Kubernetes
2. Demonstrate the I/O capabilities of MongoDB and Diamanti Kubernetes Platform under different workload scenarios

Specifications

For these tests we used 2 Diamanti D20 clusters each with 3 nodes, specifications described below, combined with MongoDB Kubernetes Operator and MongoDB Cloud Manager to showcase the benefits of a purpose-built platform. While we chose the Cloud Manager for this test case, MongoDB Ops Manager would be typically used for production scenarios. The entire example from the functional and Kubernetes perspective is identical using either cluster management solution.

Hardware Used

[Diamanti D20 series](#) of modern hyper-converged platforms are packaged with [Diamanti Ultima](#) I/O acceleration cards with varying configurations of Intel CPUs, memory, and NVMe storage. Diamanti Ultima is a pair of second generation PCIe based [I/O acceleration](#) cards that offload networking and storage traffic to deliver dramatically improved performance. They provide enterprise-class [storage services](#) for disaster recovery (DR) and data protection (DP) of mission critical applications, including mirroring, snapshots, multi-cluster asynchronous replication, and backup and recovery.

For these tests, we used the Diamanti D20 “Small” nodes to form the two 3 node clusters. The D20 “Small” has 20 physical CPU cores running at 2.2GHz each, 192 GB of RAM and Diamanti Ultima I/O acceleration for 40G networking, and 4 TB of ultra-fast, low latency NVMe Storage.

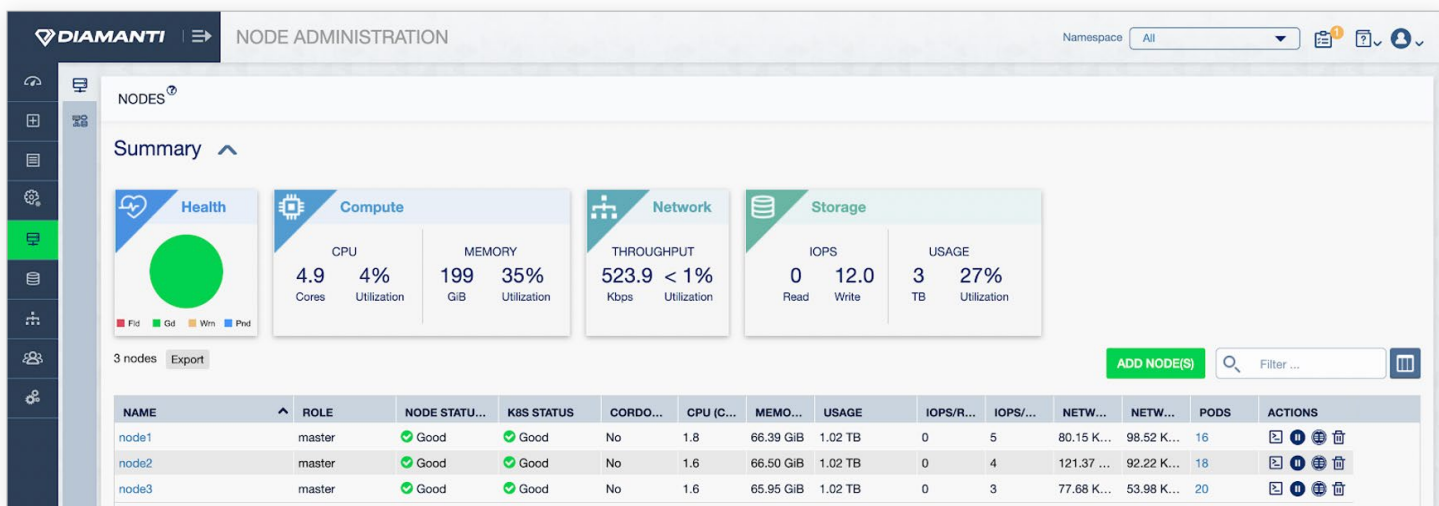


Figure 1: Easy to Manage and Monitor Diamanti D20 nodes

Software Used

Diamanti Spektra is the prevalidated, pre-packaged, and fully-featured software stack including CNCF certified Kubernetes, container runtime, operating system, enterprise-class DP/DR features, access controls, and Management UI. Our recent release of [Diamanti Spektra 2.4](#) offers protection for data at rest with AES 256-bit [volume encryption](#). With this capability, enterprise customers can prevent data theft and protect against security breaches while not incurring significant performance penalties.

The MongoDB database-as-a-service tier is implemented with MongoDB Ops Manager in conjunction with the MongoDB Enterprise Kubernetes Operator. [MongoDB Ops Manager](#) is an enterprise database management system. Ops Manager helps you automate database administration tasks such as deployment, upgrades, scaling events, and more through the platform UI and API. The [MongoDB Kubernetes Operator](#) introduces native custom resource definitions (CRDs) into your Kubernetes cluster enabling you to dynamically provision MongoDB clusters and MongoDB Ops Manager instances.

[Yahoo! Cloud Serving Benchmark](#) (YCSB) was utilized as a simple off-the-shelf synthetic test harness tool. While this tool connects to MongoDB just like any other application (with a secure connection string) we did optimize the

software specifically for this project. As such, you may not get the same performance results and neither these tests nor results are valid for any other database which YCSB may happen to support. But you will be able to use the exact same infrastructure and platform tooling for any MongoDB application running on Kubernetes.

You can try out these demonstrations yourself with the example deployments and charts available in the [total-cluster](#) Github repository.

Test Results

There are various levels of 'sophistication' you can choose for your total-cluster. The easiest way to get started is with the extra-lite option available in the setups¹. This option is "lightweight" in the sense that you'll be leveraging MongoDB Cloud Manager which is a hosted version of MongoDB Ops Manager. Using Cloud Manager perfectly positions your organization to incorporate fully hosted databases running on MongoDB Atlas, while still allowing you to utilize Diamanti hardware for your applications. More advanced configurations like this are possible through a number of cloud-native MongoDB options such as the MongoDB Atlas Service Broker, the MongoDB Atlas Terraform Provider, and the MongoDB Atlas CloudFormation Custom Resources.

1. <https://github.com/jasonmimick/total-cluster/blob/master/README.md#setups>

The first step is to head over to cloud.mongodb.com. Here you need to:

- Create an account (if you do not have one already)
- Create a Cloud Manager organization²- note to connect databases running in your own Kubernetes cluster it must be a Cloud Manager project

Features	<input type="radio"/> MongoDB Atlas	<input checked="" type="radio"/> Cloud Manager
Automated database configuration	✓	✓
Continuous backup and point-in-time recovery	✓	✓
Queryable backup snapshots	✓	✓
Fine grained database monitoring & customizable alerts	✓	✓

Figure 2: Create organization in the MongoDB Cloud Manager

- Add an Organization API Key³ as shown in Figure 3 below:

Public Key: OPITFAJV

Description*: demo key (max 250 characters)

Organization Permissions*: Organization Member

Figure 3: Create an API Key in the organization

2. <https://docs.cloudmanager.mongodb.com/tutorial/manage-organizations/#create-an-organization>

3. <https://docs.atlas.mongodb.com/configure-api-access/#manage-programmatic-access-to-an-organization>

You can find all the details for installing this in your own clusters in the [Diamanti-MongoDB YCSB Tests How to](#) document over in our repository.

Below you can see the simple deployment of MongoDB on the Diamanti Spektra management console. For all testing, we utilized a **3-way ReplicaSet using 4 GB of RAM and 4 CPU cores per container**, along with making use of the Diamanti Spektra Quality of Service (QoS) capabilities to give the containers a guarantee of 20,000 IOPS and 500 MB/s of network throughput. The setup is shown below.

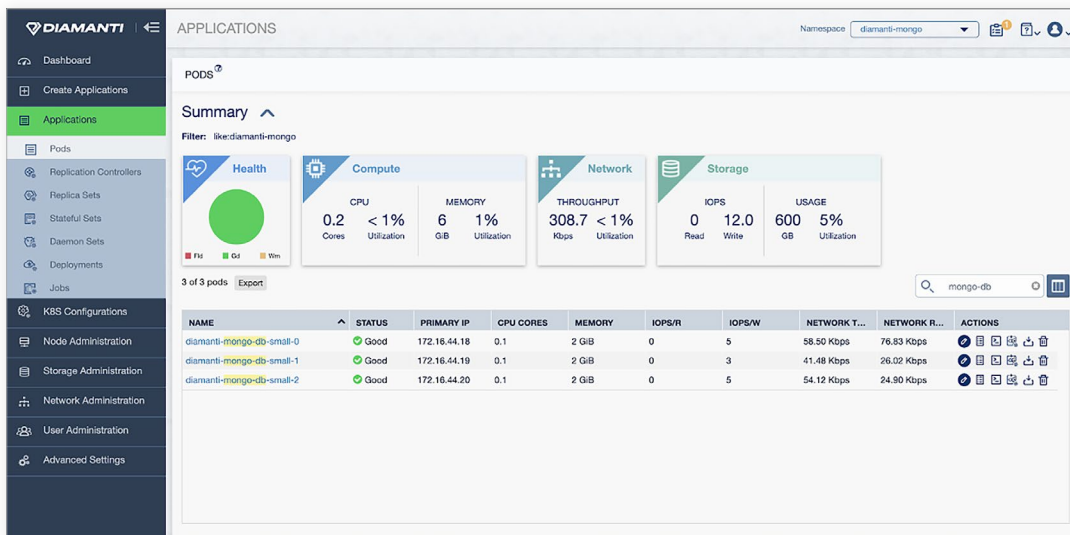


Figure 4: Diamanti MongoDB 3-way ReplicaSet overview

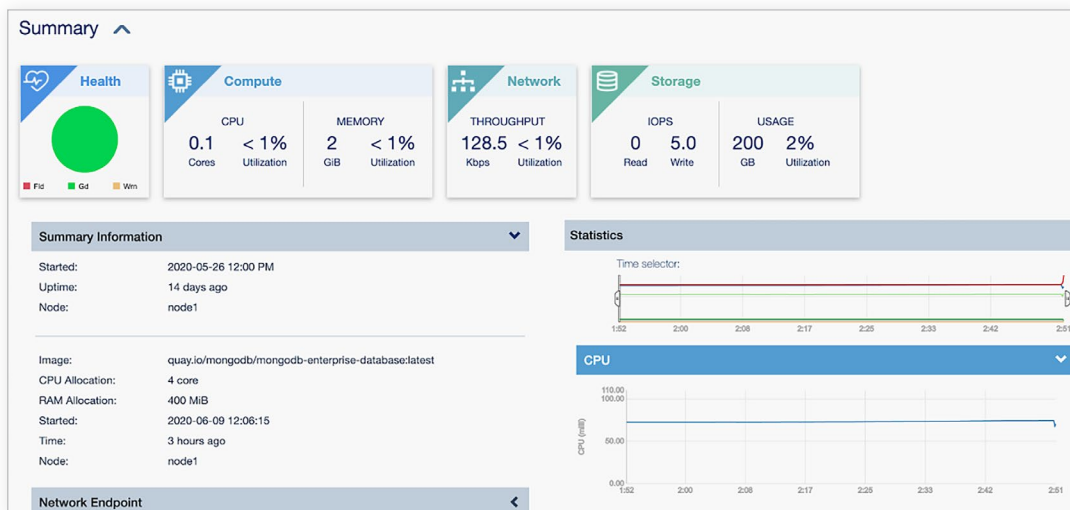


Figure 5: Diamanti container level metrics

Figure 4 shows the Diamanti Spektra dashboard displaying the 3 deployed containers in the namespace **diamanti-mongo**. Figure 5 shows the detailed metrics of one of the deployed containers.

One of the built-in features of the Diamanti platform is its ability to augment Kubernetes by providing a way for users to define QoS tiers, the default set of which is shown in the figure below.

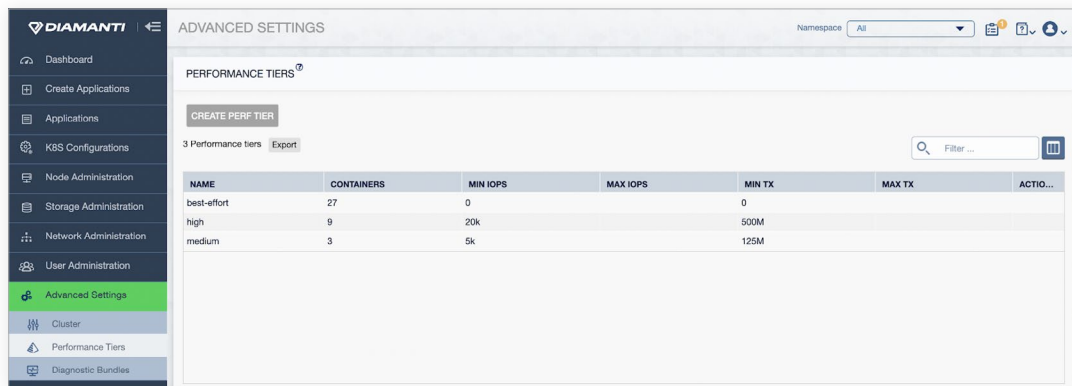


Figure 6: Diamanti default set of performance tiers

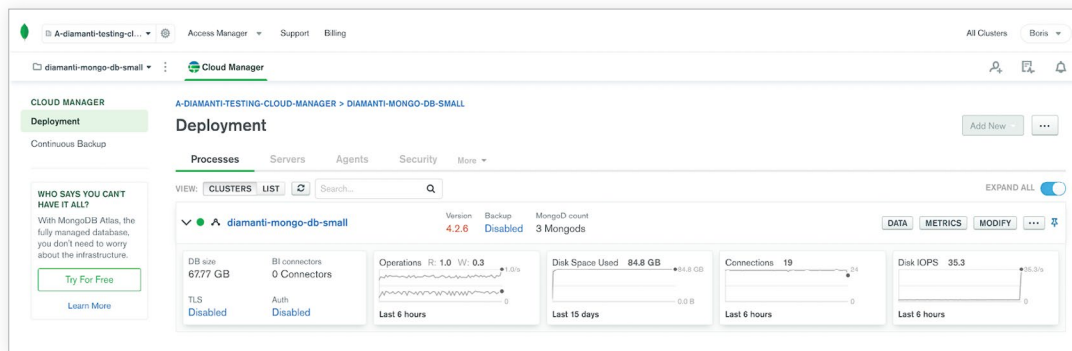


Figure 7: CloudManager view of the deployed cluster

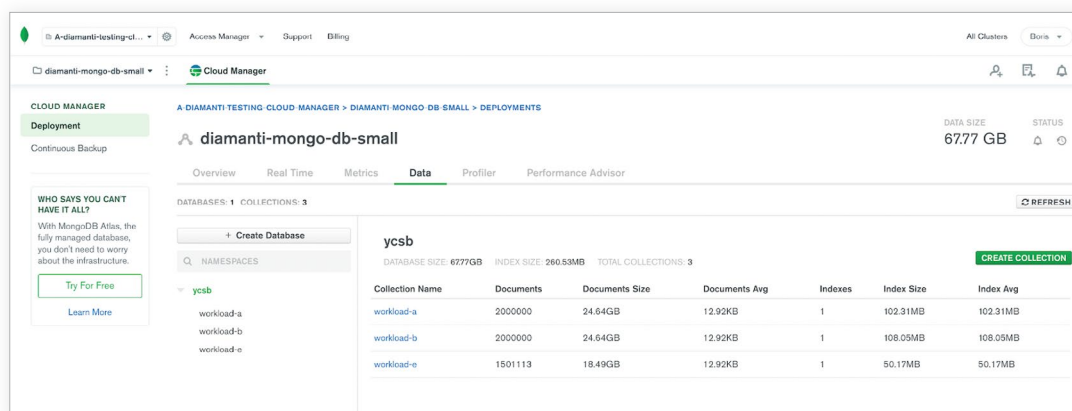


Figure 8: CloudManager view of the loaded data for the experiments

Tests I - Workload A (50/50 read/write)

The overall strategy for all of the completed tests was to use as much of the [default workload](#) configurations as possible. The notable changes are highlighted for each test. The most notable of these changes was the use of much larger **fieldcount** and **fieldlength** parameters than what is used by default. This was done in order to get the tests to behave closer to real-world usages of MongoDB.

We would like to note that these numbers are chosen specifically for approximating real world usages, however, each application and usage can have vastly different sizing requirements which will produce different outcomes.

```
recordcount=1000000      scanproportion=0
operationcount=10000000  insertproportion=0
workload=site.ycsb.workloads.CoreWorkload requestdistribution=zipfian
readallfields=true       table=workload-a
readproportion=0.5       fieldcount=50
updateproportion=0.5     fieldlength=250
                        threadcount=16
```

Tests II - Workload B (95/5 read/write)

```
recordcount=1000000      scanproportion=0
operationcount=10000000  insertproportion=0
workload=site.ycsb.workloads.CoreWorkload requestdistribution=zipfian
readallfields=true       table=workload-b
readproportion=0.95it    fieldcount=50
updateproportion=0.05    fieldlength=250
                        threadcount=16
```

Tests III - Workload E (95/5 Scan/Update)

```
recordcount=1000000      requestdistribution=zipfian
operationcount=10000000  maxscanlength=100
workload=site.ycsb.workloads.CoreWorkload scanlengthdistribution=uniform
readallfields=true       table=workload-e
readproportion=0         fieldcount=50
updateproportion=0       fieldlength=250
scanproportion=0.95      threadcount=16
insertproportion=0.05
```


Findings

First and foremost, we would like to note that we are not providing outputs of the LOAD phase for these tests, as they are inconsequential and their sole function is to pre-load data to the cluster. The specifics of the runs can be seen in Figure 8 showing each individual table created for the runs we are executing, as well as individual table statistics and information.

Test I - Workload A

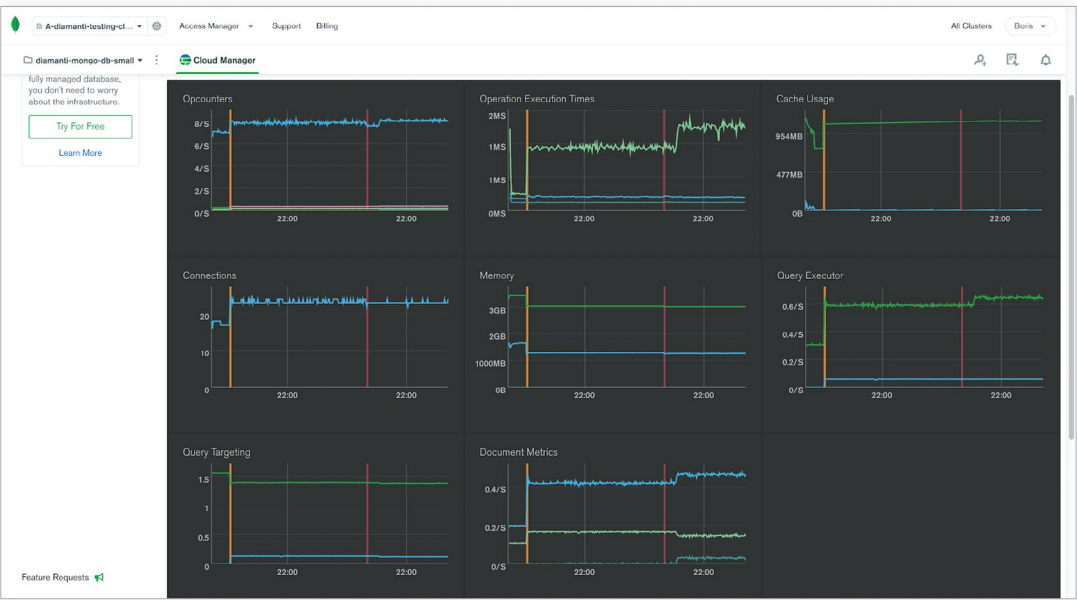


Figure 9: CloudManager view during Workload A run



Figure10: CloudManager Real-Time statistics view during the run

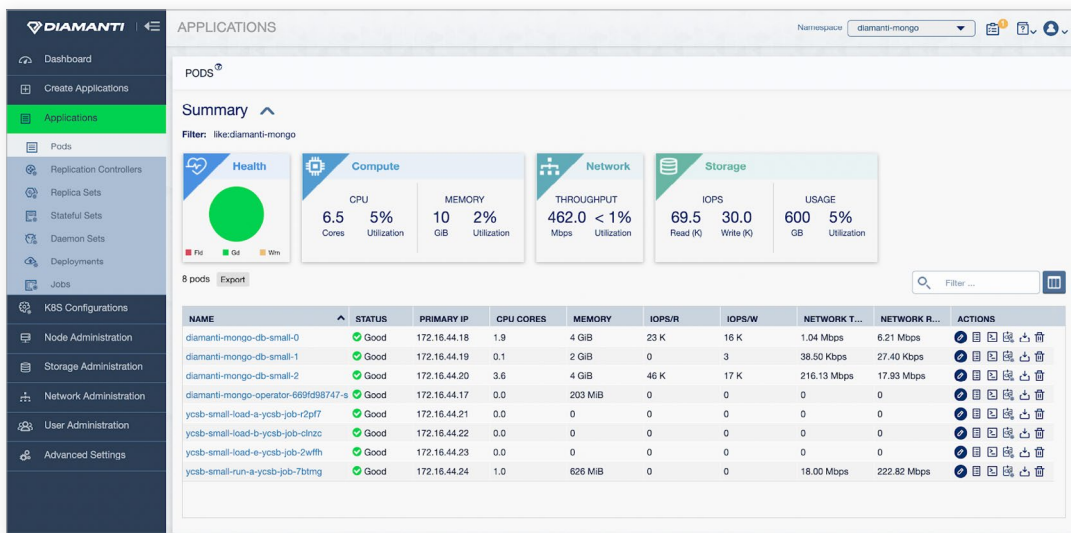


Figure11: Diamanti view during the run phase

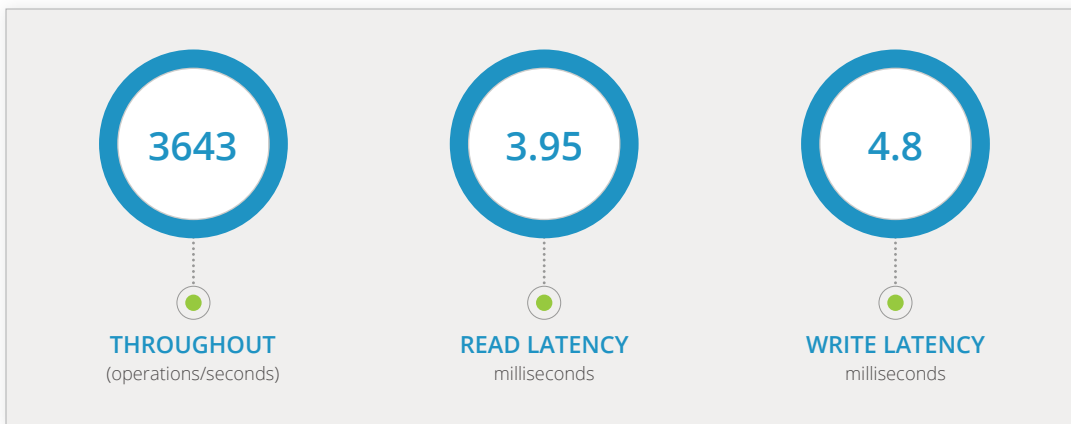


Figure12: Summary of the test run for Workload A

Raw Numbers for Workload A:

```
[OVERALL],RunTime(ms),2744332
[OVERALL],Throughput(ops/sec),
3643.8739919222603
[READ],Operations,4998747
[READ],AverageLatency(us),
3956.550892653699
[READ],MinLatency(us),280
[READ],MaxLatency(us),392959
[READ],95thPercentileLatency(us),2313
[READ],99thPercentileLatency(us),75327
[CLEANUP],Operations,16
[CLEANUP],AverageLatency(us),842.937
[CLEANUP],MinLatency(us),2
[CLEANUP],MaxLatency(us),13247
[CLEANUP],95thPercentileLatency(us),204
[CLEANUP],
99thPercentileLatency(us),13247
[UPDATE],Operations,5001253
[UPDATE],AverageLatency(us),
4810.487657792957
[UPDATE],MinLatency(us),283
[UPDATE],MaxLatency(us),393727
[UPDATE],95thPercentileLatency(us),4959
[UPDATE],99thPercentileLatency(us),75903
```

Test II - Workload B

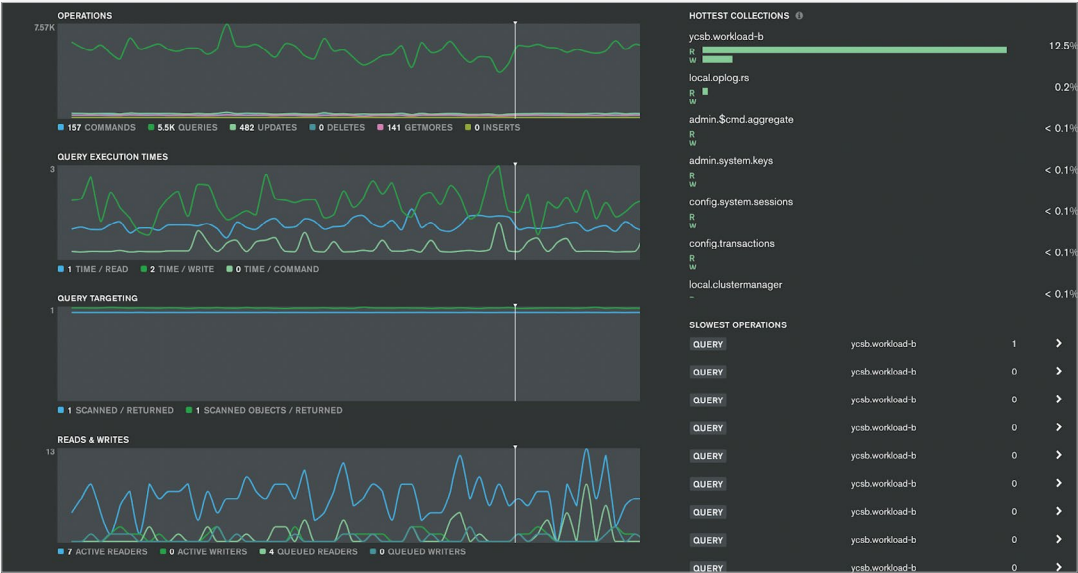


Figure13: CloudManager Real-Time statistics view during the run

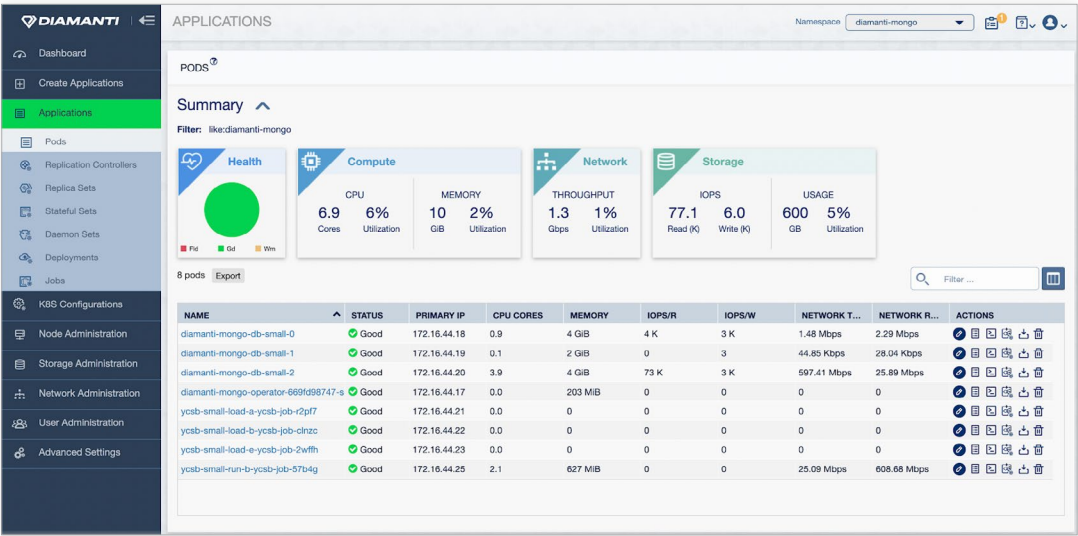


Figure14: Diamanti view during the run

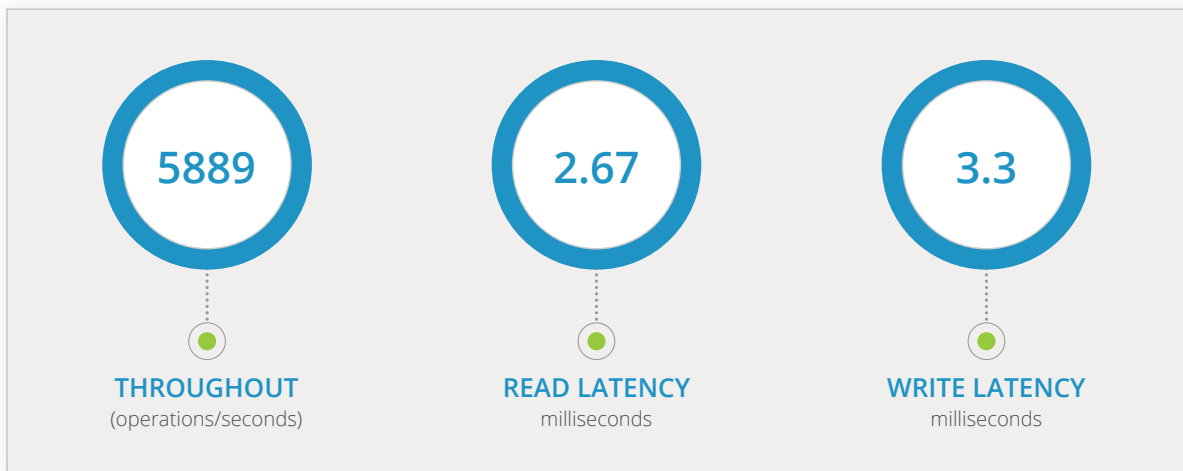


Figure15: Summary of the test run for Workload B

Raw Numbers for Workload B:

```
[OVERALL],RunTime(ms),1697923
[OVERALL],Throughput(ops/sec),5889.5485837697
[READ],Operations,9500301
[READ],AverageLatency(us),2678.458995772871
[READ],MinLatency(us),222
[READ],MaxLatency(us),133503
[READ],95thPercentileLatency(us),1860
[READ],99thPercentileLatency(us),61791
[CLEANUP],Operations,16
[CLEANUP],AverageLatency(us),790.75
[CLEANUP],MinLatency(us),2
[CLEANUP],MaxLatency(us),12591
[CLEANUP],95thPercentileLatency(us),27
[CLEANUP],99thPercentileLatency(us),12591
[UPDATE],Operations,499699
[UPDATE],AverageLatency(us),3307.282377991551
[UPDATE],MinLatency(us),319
[UPDATE],MaxLatency(us),116031
[UPDATE],95thPercentileLatency(us),2241
[UPDATE],99thPercentileLatency(us),63487
```

Test III - Workload E

The starkest result here is that the scan heavy workload is producing much lower operations per second over the span of the test. This makes sense when you take into account what is happening in this test. In Figures 17 and 18, you can see all the YCSB scans are using the index and then returning the matching documents. So, if you take this plus the increase in document size - it means there are tons more I/O going on - not just one document over the wire but n-documents per query. This results in more I/O work on the backend. The important thing to note here is the fact that a 4 GB + 4 CPU core database is able to push 150K+ IOPS and do the work of a much larger instance in a substantially smaller footprint.

Note that we skip showing the CloudManager Real-time screens for this run as the database throughput and work is negligible. Instead, we are providing the CloudManager Profiler view as shown in Figures 17 and 18, which gives a better explanation of the scan operations executed by the workload. In this view you can see that a large swath of documents gets scanned, producing the backend IOPS drive, and then returned, leading to high network throughput. But since each such scan takes over 100ms and returns an arbitrary number of documents, the actual transactions per second (TPS) are pretty low.

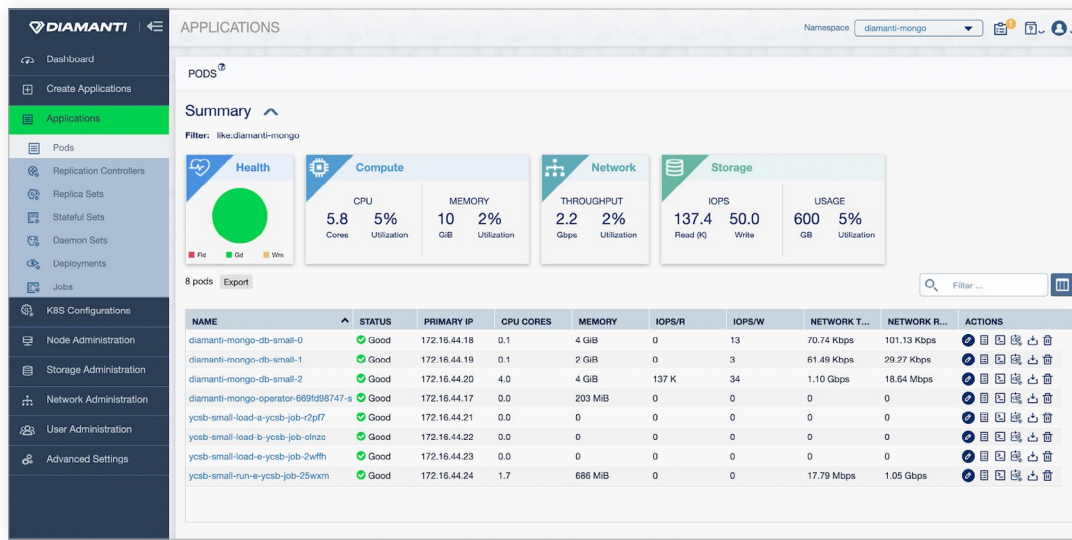


Figure16: Diamanti view during the Workload E run

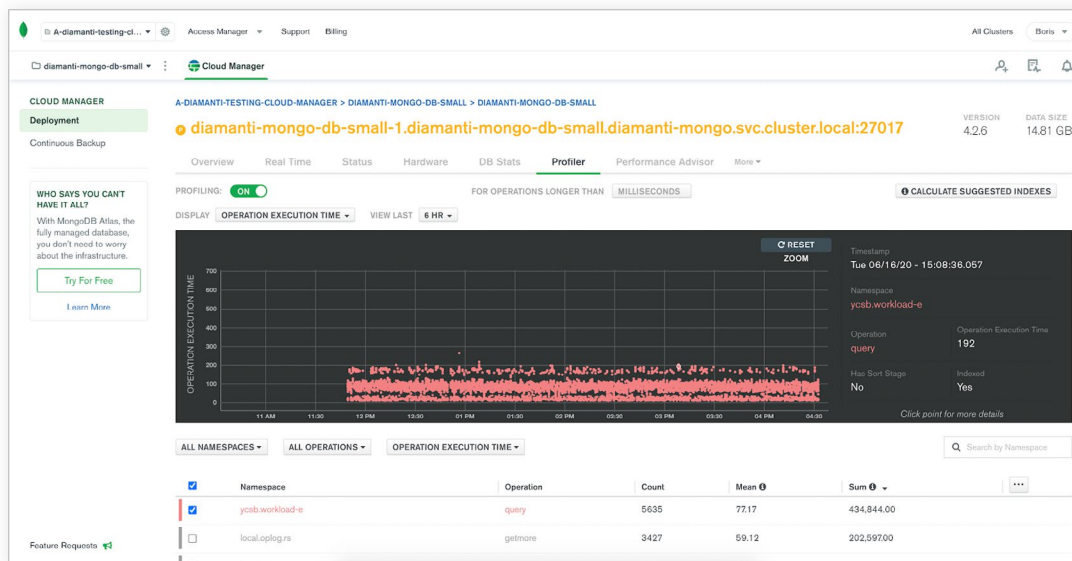


Figure17: CloudManager Profiler view, showing slow queries

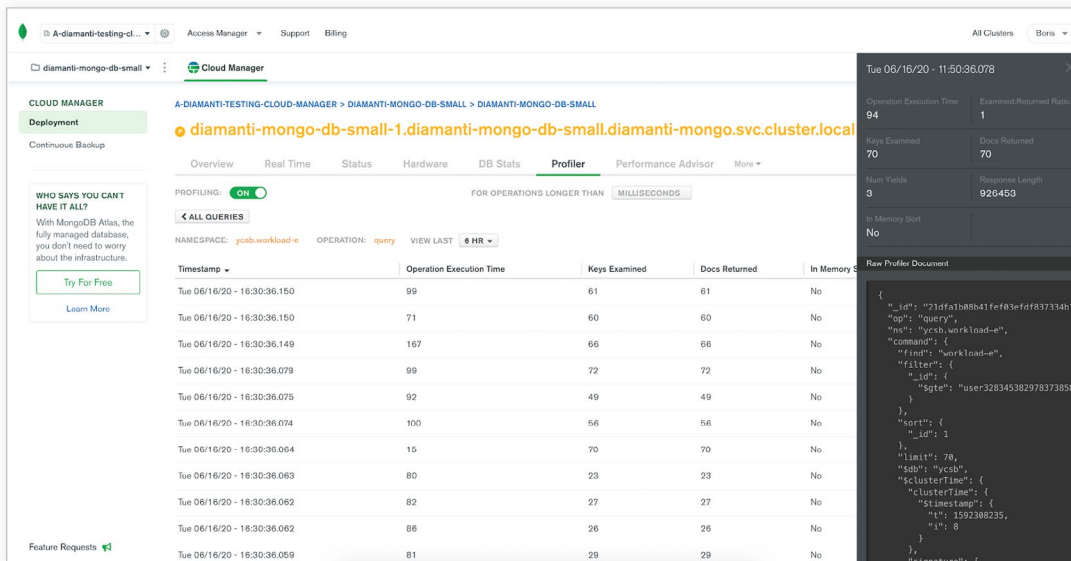


Figure 18: CloudManager Profiler drill down view of a sample slow query

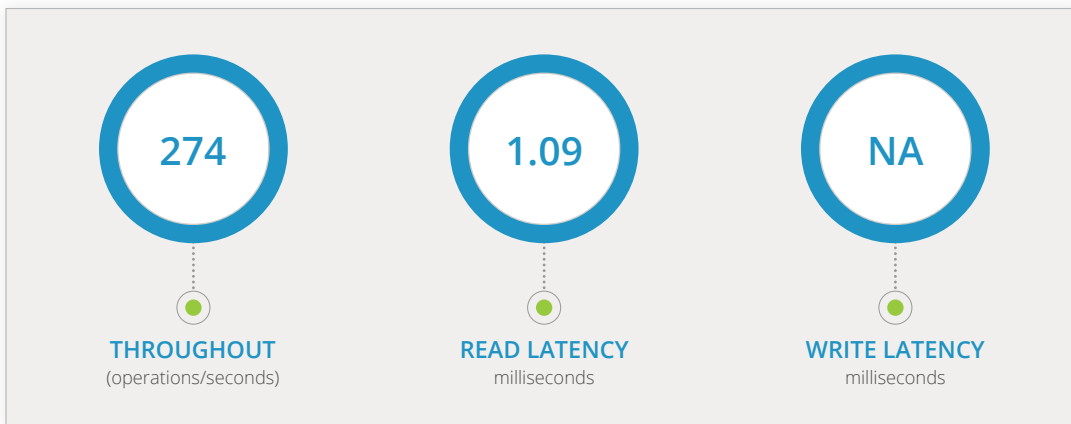


Figure 19: Summary of the test run for Workload E

Raw Numbers for Workload E:

[OVERALL],RunTime(ms),36469749	[INSERT],MinLatency(us),343
[OVERALL],Throughput(ops/sec),274.19985807963747	[INSERT],MaxLatency(us),198527
[CLEANUP],Operations,16	[INSERT],95thPercentileLatency(us),50463
[CLEANUP],AverageLatency(us),1096.5	[INSERT],99thPercentileLatency(us),60447
[CLEANUP],MinLatency(us),2	[SCAN],Operations9498887
[CLEANUP],MaxLatency(us),17471	[SCAN],AverageLatency(us),60923.269726547966
[CLEANUP],95thPercentileLatency(us),33	[SCAN],MinLatency(us),270
[CLEANUP],99thPercentileLatency(us),17471	[SCAN],MaxLatency(us),449791
[INSERT],Operations,501113	[SCAN],95thPercentileLatency(us),122751
[INSERT],AverageLatency(us),8874.33323022951	[SCAN],99thPercentileLatency(us),186879

Conclusion

Enterprises are beginning to adopt containers en masse, solidifying containers as the computing foundation for the next generation. Although Kubernetes is a powerful and modern platform for containerized workloads, it is full of complexity and difficult to deploy. Moreover, it is not a solution in and of itself, it requires many supporting elements to become a complete and operational platform fit for production use. Enterprises considering Kubernetes should look to remove as much complexity as possible through pre-built and commercially supported container platforms.

Early adopters used Kubernetes predominantly for stateless workloads due to its agility and simplicity. Stateful databases, such as MongoDB, are also increasingly being supported in Kubernetes, but require persistent storage and high performance I/O systems. While containers and Kubernetes can abstract away much of the infrastructure, the underlying physical I/O systems are still important, particularly for I/O heavy workloads like MongoDB. Thus, the underlying infrastructure is still an important consideration when containerizing high-performance workloads such as databases. Enterprises should also look to leverage Kubernetes operators, which will make provisioning and management of complex, stateful workloads easier.

What we have shown is that combining Kubernetes and MongoDB is now quite a valid and easy task, however, what is not so easy is the ability to drive performance

out of it without having a good underlying infrastructure. We have proven through the results for Workload A and B that Diamanti is capable of delivering a big punch in a tiny package. This means you can increase density for many and more complex workloads while needing to utilize less infrastructure.

The interesting results we see from Workload E also follow the same paradigm, however, it is harder to illustrate as the coveted transactions per second performance is not able to tell the full story here. Thus, we have to dive into understanding that this workload's nature and overall idea is to scan through data, which drives high backend I/O, while only returning matching documents for that transaction. This back and forth is expensive as far as I/O is concerned (this can be seen in Figure 16). While the overall TPS was marginal, we can clearly see that the same 4 GB + 4 CPU core database is able to utilize the Diamanti Ultima advantage by processing the I/O at superfast speed.

Download the IDC lab validation report [here](#) to learn more about the key findings, including:

- A comparison of operations per second between a MongoDB cluster provisioned in the public cloud and on a Diamanti cluster
- The throughput and latency observed when running on the Diamanti platform
- The simplicity of managing MongoDB in a hybrid cloud with Diamanti

